



FIT9134

Computer architecture and operating systems

Weeks 3 & 4

***Operating System II:
File Management***

File Management

File management systems allow users to store information in fundamental units called 'files'. What the file actually represents is defined by the system and/or the user.

Basic file management is typically provided by O/S, while additional file management functions may be performed by specialised software such as DBMS.

File system provides *connection between logical file structure and physical implementation*, creating **logical view** for user, and hiding **physical view**.

We will use Unix as the case study for file management in an operating system.

File management in Unix

- The Unix kernel includes the code to manage files.
- Unix utilities provide high-level file management functionality (eg. copy, delete, move, etc) to the users.
- Applications can also access the low-level file management services directly, through system calls (e.g. open, read, write etc).

for general
users

for
programmers

File Management in Unix

- O/S maintains a **directory structure** for each device, to facilitate location and organization of user files, and keeps track of free space, allocating space and reclaiming it as required.
- Provides naming, access/manipulation/storage, security/protection functions for files and directories.
- File system interacts with I/O subsystem to pass on requests for operations on storage devices

Files

- Like any operating system, Unix works on the concept of files, which are stored in a *file system*. But the concept of a file in Unix is more profound than in other operating systems :

“Everything is a file”

- In Unix, files are simply a collection of bytes stored on the storage medium. They can represent any of the following types:
 - **Ordinary Files**
 - Data (e.g. a text file, program source code)
 - Executables (e.g. a Unix command, a shell script, etc)
 - **Directories**
 - A directory is another type of file in Unix - a special "**file**" that can contain other files and other directories.
 - **Special Files**
 - Other types of files, eg. files that represent hardware devices like hard drives.

Files – naming conventions

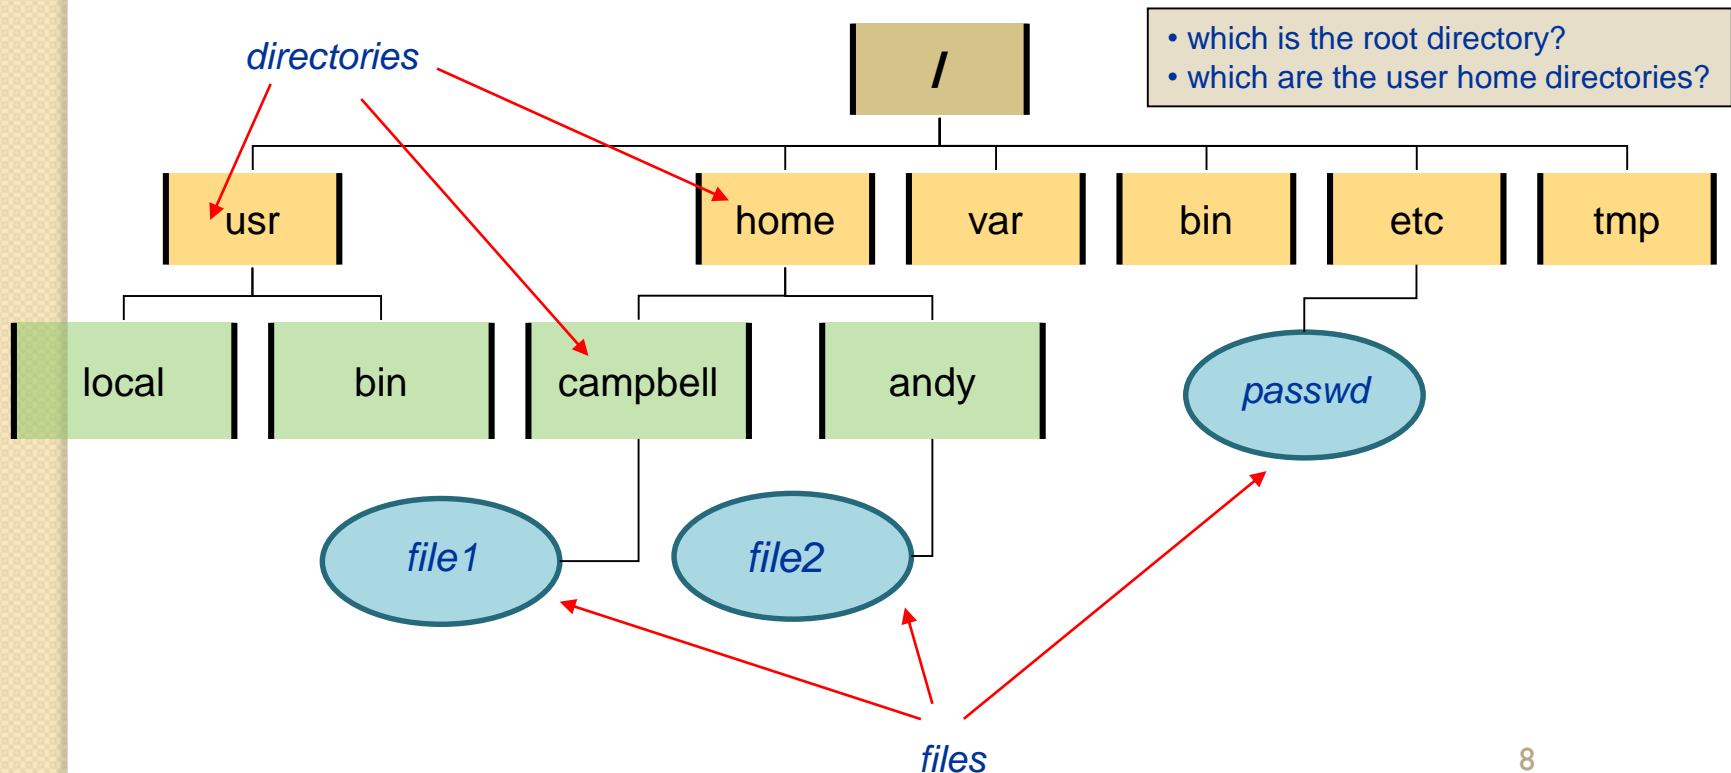
- Unix is **case-sensitive**. In general, **most Unix commands are in lowercase letters** (ie. Unix does not like uppercase letters!).
- Unix filenames are generally made up of lowercase and uppercase letters, digits, dots (.) and commas. Using spaces (or other “special” characters) in filenames can occasionally make file-handling difficult, so try to avoid them if possible.
- There is no notion of a file “extension” in Unix (unlike O/S such as Microsoft Windows). While files can have an extension, the extension (ie. the bit after a "dot") has no special meaning, and does not necessarily define the type of the file or indicate how it should be dealt with by an application.

Files - naming conventions

- In theory, any name can be used for a file or a directory (with the exception of the “**root directory**” which must always be named **/**). Unix is also very generous with the length of a filename.
- Try to avoid using special characters.
- However, important system files and directories are generally given the same “standard” names on Unix systems. Examples of some common/typical **system directories** include (but are not limited to) **/etc**, **/bin**, **/home**, **/mnt**, **/usr**, **/var**, **/tmp**, **/proc** and **/lib**.
 - you should not modify/delete these system directories/files unless you know what you are doing!

The UNIX File System

- UNIX stores files on the disk in a **hierarchical** structure.
- The top of the hierarchy is referred to as the **root directory** and is always named **/**
- Eg. a typical Unix file system might look like :



The Working Directory

- When working on a UNIX system, you are always working within a particular directory in the file system. This is called either the *working directory* or the *current directory*.
- When you first log in to the system, your working directory is set to a directory that is unique to you as a user; your *home directory*. The home directory is the directory which belongs to you and will contain your files (and you have full control over these files).

Paths (or Pathnames)

- A “path” or “pathname” represents the position/location of a file (or directory) in the file system hierarchy. Directories in the path are separated by the `/` symbol,
 - Eg. `/home/andy/week1`, `/etc/passwd`, `../week2`
- Special directory names used in paths:
 - `.` = **current** directory
 - `..` = **parent** directory (the directory immediately above the current directory in the hierarchy)
 - `~` = the user’s **home** directory

Absolute and Relative Paths

- A path that begins with **/** is called an *absolute path* and indicates the “absolute” position of the file regardless of the directory in which the user is currently working, eg. **/home/andy/letter**, **/home/david/week3**
- A path that starts from the current directory is called a *relative path* and indicates the position of the file **relative** to the directory in which the user is currently working, eg. **week3/exerciseQ2**, **../myfolder/test/mydocument**, **../../week5**
- We generally regard a path like **~/week3** to be an absolute path (even if it does not start with a **/**). Can you see why?

Ownership of files

- Files are “owned” by a user. We can check a file’s ownership using the **ls** command with the long format, eg :

\$ ls -l

← *Note : the ‘\$’ (the command prompt) is **not** part of the command!*

drwxr-xr-x	18	cheng	fit9134	2048 Nov 12 2013	WWW
drwx-----	3	cheng	fit9134	4096 Sep 26 2013	mail
-rw-r--r--	1	cheng	fit9134	1313 May 7 13:15	letter1
drwxr-xr-x	5	cheng	fit9134	2048 Aug 16 2013	tmp
drwxr-xr-x	4	cheng	fit9134	2048 Aug 5 13:48	week6

** more about this concept later **

A special file : `/etc/passwd`

- On Unix systems, there is a special file named “`passwd`” in the directory `/etc`. This file holds important information about all the `users` in the system. For instance, you may find in this file entries such as :

```
cheng:x:500:700:Andy Cheng:/home/cheng:/bin/tcsh  
cwilson:x:501:701:Campbell Wilson:/home/cwilson:/bin/bash
```

- There is also a similar file, `/etc/group`, which holds information about all the `groups` in the system.

More on Files

- File system is “flat” under Unix kernel (i.e. there is no real distinction between directories and files under Unix kernel).
 - a device is represented as a file.
 - a program is a file.
 - a directory is really also a file (!)
- however, in the user interface, the user is typically presented with a hierarchical view (eg. the folders as shown in a graphical File Manager)

More on Files

- The kernel does not identify files by names; it uses a unique number to identify a file called the *inode number*.
- The **stat** command shows detailed info about files, eg:


```
$ stat /home/cheng/recipe
  File: `/home/andy/recipe'
  Size: 11          Blocks: 8          IO Block: 4096   regular file
Device: 841h/2113d Inode: 556730       Links: 1
Access: (0664/-rw-rw-r--)  Uid: ( 1000/   andy)   Gid: ( 1000/   andy)
Access: 2015-08-09 19:32:27.875891467 +1000
Modify: 2015-08-09 19:32:27.875891467 +1000
Change: 2015-08-09 19:32:27.875891467 +1000
```

Using Unix Commands

- The original Unix O/S only allows users to interact with it via text commands (ie. user types in commands via a terminal). This is known as the “**Command-Line**” interface.
- Modern day Unix usually also provides a **Graphical** interface, allowing user to interact via pointing-and-clicking using devices such as a mouse. This makes it easier for the casual user to interact with the O/S.

Graphical or Command-Line Interface???

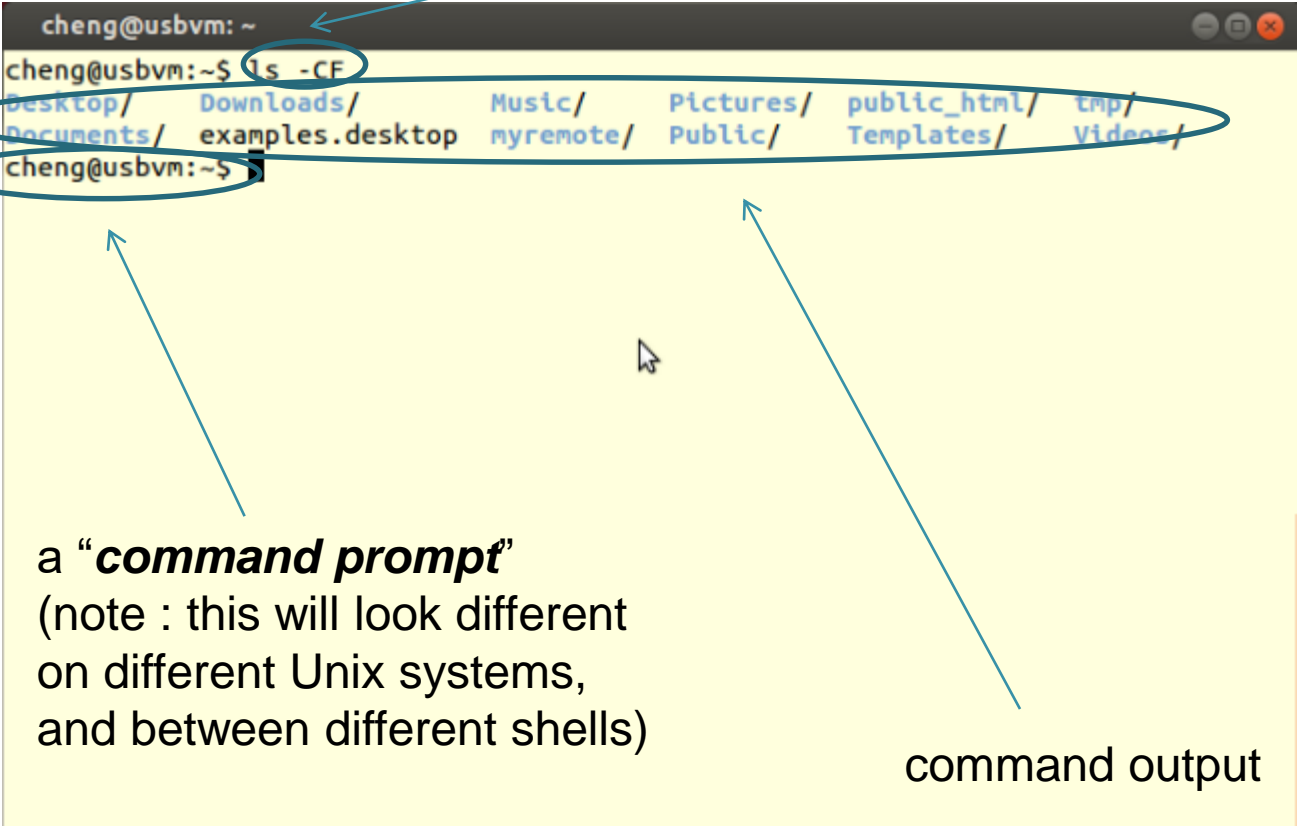
- The command-line interface is still popular among advanced Unix users, because it is generally faster & more flexible.
 - A “*Terminal Window*” is now typically provided for this purpose, eg. in Ubuntu, the user can activate that window via the Dash Home menu or shortcut-key (***ctr-alt-T***)



Command-Line Example

- In the command-line window, commands are typed in by the user, eg. :

a command
(ls -CF)



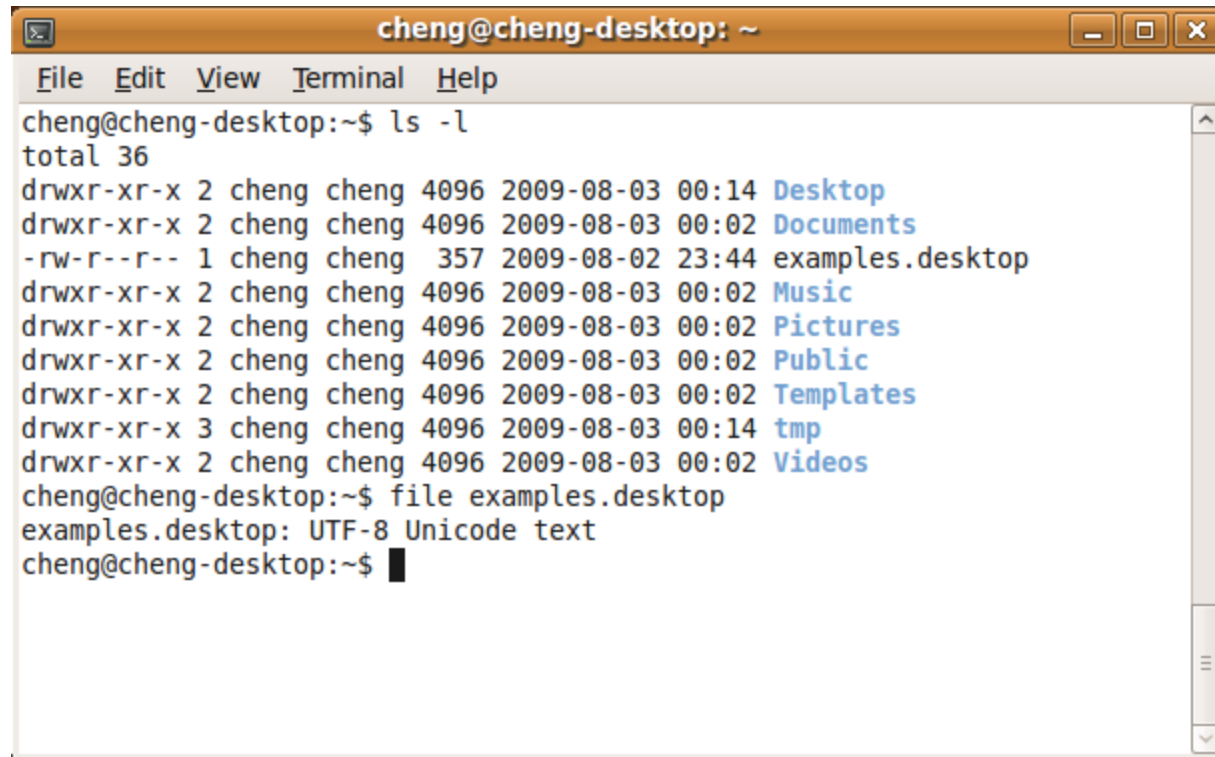
The screenshot shows a terminal window with the title bar 'cheng@usbvm: ~'. The prompt is 'cheng@usbvm:~\$'. The command 'ls -CF' is entered and circled in blue. The output is a two-line directory listing: 'Desktop/ Downloads/ Music/ Pictures/ public_html/ tmp/' and 'Documents/ examples.desktop myremote/ Public/ Templates/ Videos/'. The entire output line is circled in blue. Below the terminal, there are two blue arrows pointing to the prompt and the output. The arrow pointing to the prompt is accompanied by the text 'a "command prompt" (note : this will look different on different Unix systems, and between different shells)'. The arrow pointing to the output is accompanied by the text 'command output'.

```
cheng@usbvm: ~  
cheng@usbvm:~$ ls -CF  
Desktop/ Downloads/ Music/ Pictures/ public_html/ tmp/  
Documents/ examples.desktop myremote/ Public/ Templates/ Videos/  
cheng@usbvm:~$
```

a “**command prompt**”
(note : this will look different
on different Unix systems,
and between different shells)

command output

A simple command example



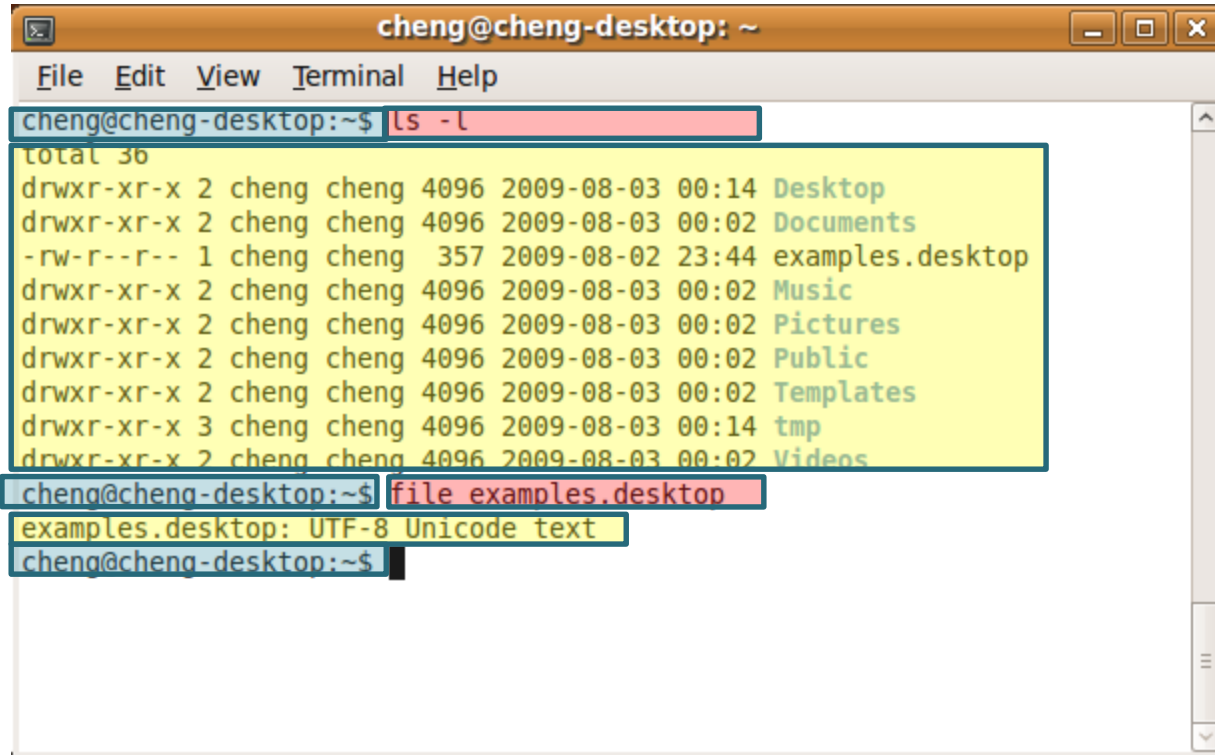
The screenshot shows a terminal window titled "cheng@cheng-desktop: ~". The window has a menu bar with "File", "Edit", "View", "Terminal", and "Help". The terminal content shows the user running the command "ls -l", which outputs a long listing of files and directories in the home directory. The files listed are Desktop, Documents, examples.desktop, Music, Pictures, Public, Templates, tmp, and Videos. The user then runs the command "file examples.desktop", which outputs "examples.desktop: UTF-8 Unicode text". The prompt "cheng@cheng-desktop:~\$" is visible at the end of the output.

```
cheng@cheng-desktop: ~  
File Edit View Terminal Help  
cheng@cheng-desktop:~$ ls -l  
total 36  
drwxr-xr-x 2 cheng cheng 4096 2009-08-03 00:14 Desktop  
drwxr-xr-x 2 cheng cheng 4096 2009-08-03 00:02 Documents  
-rw-r--r-- 1 cheng cheng 357 2009-08-02 23:44 examples.desktop  
drwxr-xr-x 2 cheng cheng 4096 2009-08-03 00:02 Music  
drwxr-xr-x 2 cheng cheng 4096 2009-08-03 00:02 Pictures  
drwxr-xr-x 2 cheng cheng 4096 2009-08-03 00:02 Public  
drwxr-xr-x 2 cheng cheng 4096 2009-08-03 00:02 Templates  
drwxr-xr-x 3 cheng cheng 4096 2009-08-03 00:14 tmp  
drwxr-xr-x 2 cheng cheng 4096 2009-08-03 00:02 Videos  
cheng@cheng-desktop:~$ file examples.desktop  
examples.desktop: UTF-8 Unicode text  
cheng@cheng-desktop:~$
```

In the diagram above, can you identify examples of these?

- a *command*
- a *command output*
- a *command prompt*

A simple command example



The screenshot shows a terminal window titled "cheng@cheng-desktop: ~". The menu bar includes "File", "Edit", "View", "Terminal", and "Help". The prompt is "cheng@cheng-desktop:~\$". The command "ls -l" is entered and highlighted in pink. The output is highlighted in yellow and lists files with their permissions, sizes, owners, dates, and names. The command "file examples.desktop" is entered and highlighted in pink, with its output "examples.desktop: UTF-8 Unicode text" highlighted in yellow. The prompt "cheng@cheng-desktop:~\$" is highlighted in blue.

```
cheng@cheng-desktop:~$ ls -l
total 36
drwxr-xr-x 2 cheng cheng 4096 2009-08-03 00:14 Desktop
drwxr-xr-x 2 cheng cheng 4096 2009-08-03 00:02 Documents
-rw-r--r-- 1 cheng cheng  357 2009-08-02 23:44 examples.desktop
drwxr-xr-x 2 cheng cheng 4096 2009-08-03 00:02 Music
drwxr-xr-x 2 cheng cheng 4096 2009-08-03 00:02 Pictures
drwxr-xr-x 2 cheng cheng 4096 2009-08-03 00:02 Public
drwxr-xr-x 2 cheng cheng 4096 2009-08-03 00:02 Templates
drwxr-xr-x 3 cheng cheng 4096 2009-08-03 00:14 tmp
drwxr-xr-x 2 cheng cheng 4096 2009-08-03 00:02 Videos

cheng@cheng-desktop:~$ file examples.desktop
examples.desktop: UTF-8 Unicode text

cheng@cheng-desktop:~$
```

In the diagram above, can you identify examples of these?

- a **command**
- a **command output**
- a **command prompt**

File types – as shown by **ls -l**

Character	File type
-	regular (ordinary) file
d	directory
b	buffered special file (e.g. a disk drive)
c	unbuffered special file (e.g. a terminal)
l	symbolic link
p	pipe
s	socket

The command **ls -l** will show things such as file types, permissions, file sizes, modification dates, etc

a typical
command prompt

The **file** command shows the type of the content of the given file name :

\$ file exercise
exercise: ascii text

the actual
command

the command output

**Note : this is how a command will
typically be shown in our lecture slides**

Pre-reading for Week 4

- The next few slides (**Slides #23-28 – on file permissions and how to change them**) are important for your Week 4's lab tasks. Please make sure you study them (plus do some online research on these topics) **before** attending the labs.

File permissions

- Three levels of permissions :
 - the **user**
 - the user's **group** and
 - **others** who have account on the system
- Three kinds of permissions (for each level) :
 - **read**, **write** and **execute**
 - these have the usual meaning for ordinary files (for directory files – x means something different – more later)

Hence 9 different combinations in total

File permissions

- A total of nine (9) binary **bits** representing the permissions:

user

r/- w/- x/-

group

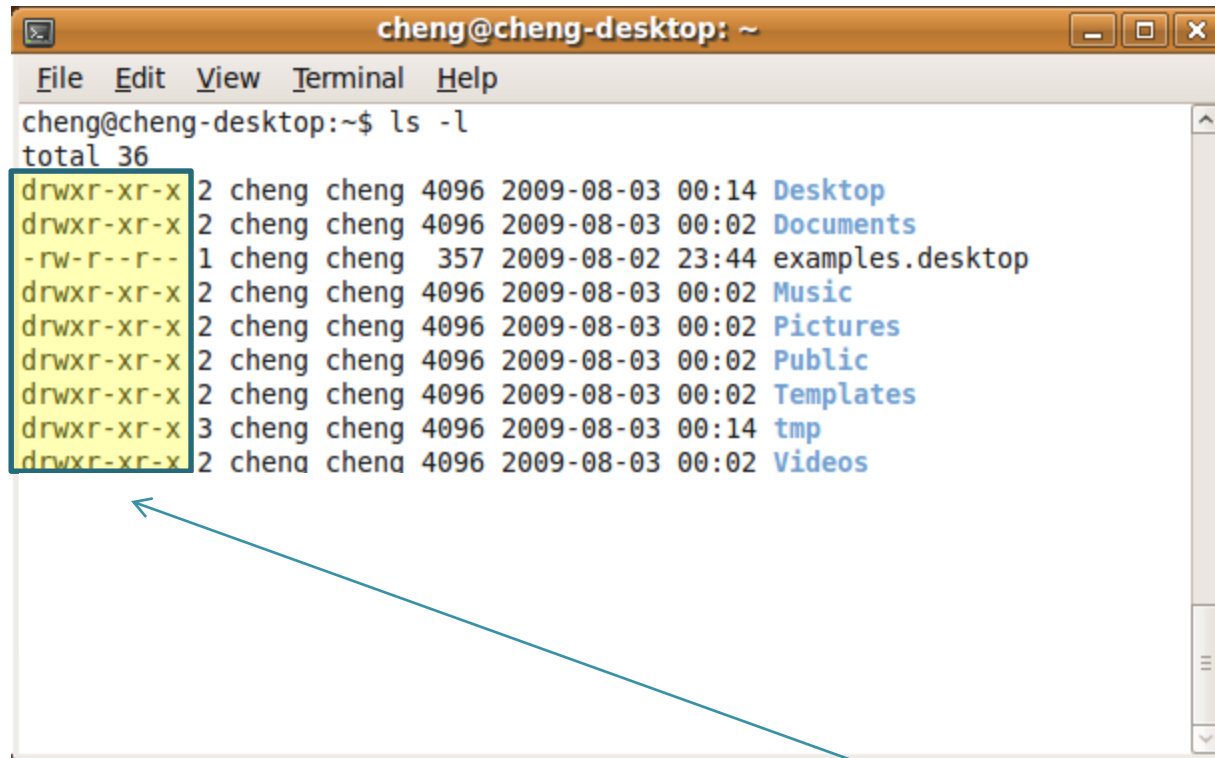
r/- w/- x/-

others

r/- w/- x/-

- a “-” indicates the permission is “off”, eg. **rw-rw-r--**
 - an example to follow shortly
- A user can choose to restrict access to his files/directories, so that other users may or may not access them.
- A **Superuser** (the “root” user) has access to all files irrespective of permissions.

Using ls -l to show file permissions



```
cheng@cheng-desktop: ~  
File Edit View Terminal Help  
cheng@cheng-desktop:~$ ls -l  
total 36  
drwxr-xr-x 2 cheng cheng 4096 2009-08-03 00:14 Desktop  
drwxr-xr-x 2 cheng cheng 4096 2009-08-03 00:02 Documents  
-rw-r--r-- 1 cheng cheng 357 2009-08-02 23:44 examples.desktop  
drwxr-xr-x 2 cheng cheng 4096 2009-08-03 00:02 Music  
drwxr-xr-x 2 cheng cheng 4096 2009-08-03 00:02 Pictures  
drwxr-xr-x 2 cheng cheng 4096 2009-08-03 00:02 Public  
drwxr-xr-x 2 cheng cheng 4096 2009-08-03 00:02 Templates  
drwxr-xr-x 3 cheng cheng 4096 2009-08-03 00:14 tmp  
drwxr-xr-x 2 cheng cheng 4096 2009-08-03 00:02 Videos
```

In the diagram above, the file permissions (marked in yellow) are shown as a **9-bit** pattern (**rw-rw-rw-**), plus a “**file-type**” bit at the front (normally “-” for file, “d” for directories)

File permissions : example using `ls -l`

```
$ ls -l examples
```

```
-rw-r--r-- 1 cheng users 357 2009-08-02 23:44 examples
```

Output explanations :

- The permission mode of this file is **read** and **write** for the **owner**, **read** only for the **group** and **read** only for **others**
- There is **1 hard link** (more on this later)
- The user-id of the file's owner is **cheng**
- The group-id of the file is **users**
- The size of the file is **357** “blocks” – NB. block size can vary between systems
- The file was last modified on **2009-08-02 23:44**
- The file name is **examples**

The option “**-l**” in the command above is to request the output in **long** format
There is another option, “**-h**”, which will make `ls` display sizes in “human readable” format (eg. 8K, 555M, 4G, etc)

Change file permission (chmod)

Syntax: `chmod [-R] who [op] [permission] file-list`

- **who** is one of:
 - **u** user owner of the file
 - **g** group group to which the owner belongs
 - **o** other all other users
 - **a** all can be used in place of u,g,o
- **op** is one of:
 - **+** add permission
 - **-** remove permission
 - **=** set permission (reset all other permissions)
- **permission** is one or more of : **r, w, x**
- Note : this is typical of a Unix command – the command is given with some option(s) & the actual operands)

chmod: Examples

You can also combine the options,
eg. `chmod u+x,o+w temp`

NB. no space here
↑

```
ls -l temp
```

```
-rw-rw-r-- 1 cheng users 57 Jul 12 16:47 temp
```

```
chmod a=rw temp
```

(set all permissions)



```
-rw-rw-rw- 1 cheng users 57 Jul 12 16:47 temp
```

```
chmod o=w temp
```

(set others' permissions)



```
-rw-rw-r-- 1 cheng users 57 Jul 12 16:47 temp
```

```
chmod u+x temp
```

(set user's permissions)



```
-rwxrw-r-- 1 cheng users 57 Jul 12 16:47 temp
```

Note how the **chmod** commands change the file's permissions

File access for processes

- When a process executes, it has four id's:
 - a **real** user-id
 - an **effective** user-id
 - a **real** group-id
 - an **effective** group-id
- these id's determine a process's access permissions to files/directories.

Real versus Effective user id's

- *Real UID* is the *UID* of the user that created *THIS* process – ie. the user who executes/runs the program.
- *Effective UID* is used to evaluate privileges of the process to perform a particular action.
- This distinction is useful, since it allows a normal user to sometimes execute processes which require higher access privileges than he has,
 - eg. a build-in program may be available for everyone to execute, but may need some super-user privileges to carry out certain system-related tasks. So we can change the **effective UID** of that process to make the O/S think that the super-user is executing that program (instead of the user who actually runs the program).

File access for processes

- When a process tries to access a file, which of the three (**User**, **Group**, or **Others**) ownership permission applies?
- The general rules are as follows:
 - If the process's effective user-id is same as the owner of the file then **User** permissions apply
 - Otherwise, if the process's effective group-id is the same as file's group id then **Group** permissions apply
 - Otherwise, **Others** permissions apply
- In other words, what a user is permitted to do to a file depends on whether he owns the file, his group owns the file, or neither...

setuid and setgid

- A process' access privileges depend on who **executes** the process, not on who **owns** the executable program itself
 - This is safer in general, but not helpful in some (rare, but important) cases.
 - This can be overcome using special permissions: *set-user-id* and *set-group-id*
 - When a program with *setuid* permission is executed, the resulting process's effective *user-id* becomes that of owner of the program (instead of the user who executes that program).
 - Similarly with *setgid*.
 - In both cases, the real *uid* and *gid* are not affected

Example of using set-uid

- */etc/passwd* file stores the encrypted passwords (plus other info) of all registered users in the system

who has access
to this file??

- `$ ls -l /etc/passwd` ← the password file

- `-rw-r--r-- 1 root root 29757 Jul 23 9:05 /etc/passwd`

- The command *passwd* (in */usr/bin*) needs to be used by any user to change his/her own password, i.e. to modify */etc/passwd*. How can this be possible when a normal user does not have write permission to the */etc/passwd* file?
 - the solution : applying setuid on the *passwd* command

Example of using set-uid

- Let's examine the permissions of the */usr/bin/passwd* utility :

```
$ ls -l /usr/bin/passwd
```

← the password utility

```
-r-s--X--X 3 root root 16384 Feb 3 15:03 /usr/bin/passwd
```

The *s* instead of the usual *x* in the permission field of the owner means:

when this command is run (by a normal user), the process will be granted the permissions corresponding to the process-owner, in this case, **root**

- root** has write access to */etc/passwd* file (see previous slide)
- this means : users cannot directly modify the */etc/passwd* file, but can change his/her own password by using the *passwd* command. This is required for obvious security reasons...

Binary, Octal & Hexadecimal notations....

- In computing, it is often convenient to express numbers in format other than decimal notations.
- Some useful readings/links :
 - <http://kb.iu.edu/data/agxz.html>
 - http://www.allaboutcircuits.com/vol_4/chpt_1/4.html
 - <http://en.wikipedia.org/wiki/Octal>
 - <http://en.wikipedia.org/wiki/Hexadecimal>
 - <http://www.tonymarston.net/php-mysql/converter.php>

Some commonly used Octal (base-8) values for file permissions

700	==>	/* owner:	<div style="border: 1px solid black; padding: 2px;">rwx</div> ----- */
400	==>	/* owner:	<div style="border: 1px solid black; padding: 2px;">r--</div> ----- */
200	==>	/* owner:	<div style="border: 1px solid black; padding: 2px;">-w-</div> ----- */
100	==>	/* owner:	<div style="border: 1px solid black; padding: 2px;">--x</div> ----- */

070	==>	/* group:	--- <div style="border: 1px solid black; padding: 2px;">rwx</div> --- */
040	==>	/* group:	--- <div style="border: 1px solid black; padding: 2px;">r--</div> --- */
020	==>	/* group:	--- <div style="border: 1px solid black; padding: 2px;">-w-</div> --- */
010	==>	/* group:	--- <div style="border: 1px solid black; padding: 2px;">--x</div> --- */

007	==>	/* others:	----- <div style="border: 1px solid black; padding: 2px;">rwx</div> */
004	==>	/* others:	----- <div style="border: 1px solid black; padding: 2px;">r--</div> */
002	==>	/* others:	----- <div style="border: 1px solid black; padding: 2px;">-w-</div> */
001	==>	/* others:	----- <div style="border: 1px solid black; padding: 2px;">--x</div> */

4000	==>	/* set user id on execution */
2000	==>	/* set group id on execution */

Values of 0-7 can be used to indicate if a particular bit is "on or "off"

This sort of "shortcut" is commonly used in Unix commands

a '1' means the corresponding bit is "on"

a '0' means the corresponding bit is "off"

Examples :

chmod 400 file1

====>

4	0	0
r--	---	---

(4 0 0 == 100 000 000)

chmod 764 file2

====>

7	6	4
rwx	rw-	r--

(7 6 4 == 111 110 100)

Standard Input, Output and Error

- Remember, in Unix, *everything is a file...*
- Every time a shell is started, 3 files are opened automatically :

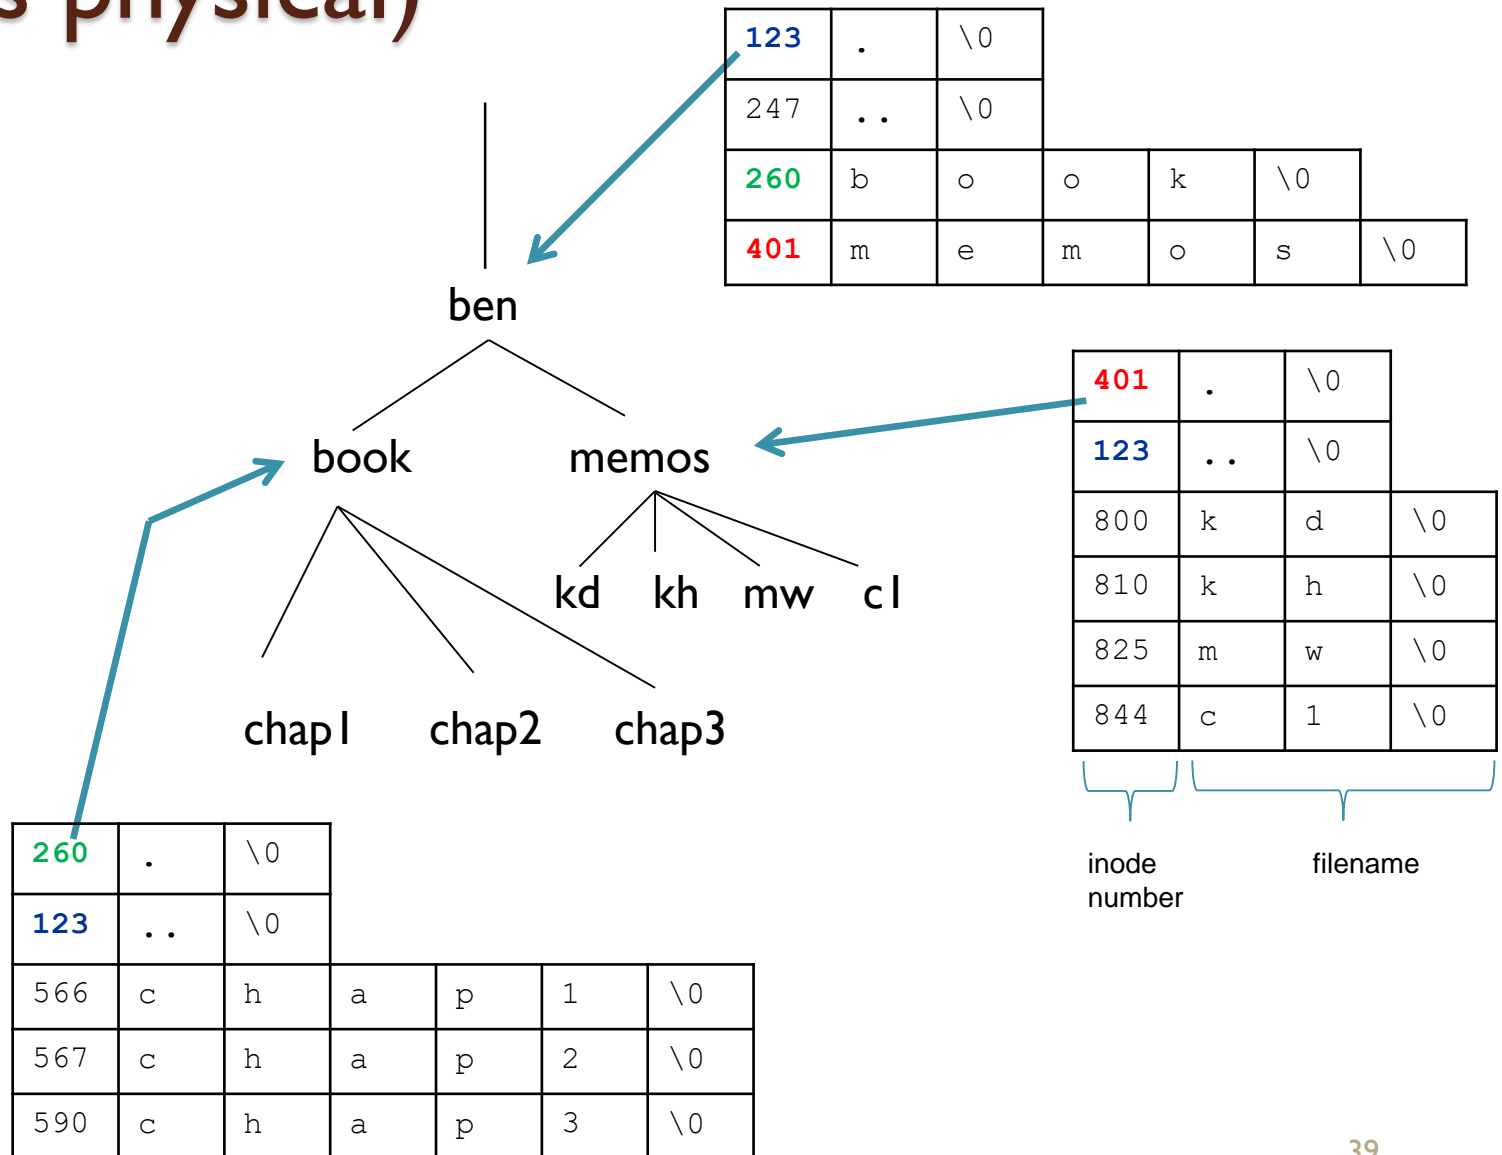
stdin, stdout, stderr

- | • File | Default Device | File Descriptor |
|---------------|-----------------------|------------------------|
| <i>stdin</i> | keyboard | 0 |
| <i>stdout</i> | screen | 1 |
| <i>stderr</i> | screen | 2 |
- A process can then easily read/write to/from these “files”; this makes I/O programming relatively easy.

Directories (Unix)

- Unix directories consist of a series of directory entries.
- Each directory entry contains at least the inode number for the file and a character field containing the filename.
- The actual format of the directory is file-system dependent – and is hidden from the user.

Example: directory structure (logical vs physical)



Directory permissions (Unix)

- **Read** permission
 - Processes can list names and subdirectories within the directory
- **Write** permission
 - Processes can alter the directory, i.e. create and remove existing files. Note: **write** permission on a file is required to modify its contents (this is not altering the directory so write permission on the directory does not say anything about this) but *not* to remove it if the directory has write permission!
- **Execute** permission
 - Allows a user to change into the directory (`cd` from the shell or `chdir` as a system call). In addition, to open a file or execute a program a user must have **execute** permission on all the directories leading to the file as specified in the file's absolute pathname.

Links – Hard Links

- A **hard** link is a pointer/reference to a file - *every file has at least one hard link to it.*
- The link is how the operating system associates a file *name* with the address of the actual data on the storage device.
- Additional links can be created to allow sharing of files or access through a different name.
- Hard links can be created in Unix using *ln*.
- *A file exists until the last hard link to it is removed.*
When the last hard link is removed, the space previously used by the file is marked for re-use.

Links – Symbolic Links

- a **symbolic** link (or **symlink**, or **soft** link) is a **file** that contains a pointer to another file.

Eg:

the 'l' indicates a soft link

name of the link

what the link is "pointing" to

```
lrwxrwxrwx 1 cheng cheng 26 2008-02-29 09:53 Examples -> /usr/share/examplecontent
```

- **symlinks can exist even if the file they pointed to does not exist!**
- symlinks must be used if a link is to span filesystems (sometimes referred to as partitions). Hard links cannot be used on directories, and do not work across different filesystems.

Unix *ln* command examples

- Hard link (default) :

ln ~/week4/myfile hardlink-1

sources

links

- Soft link (using *-s* option) :

ln -s ~/week4/myfile softlink-1

specifies a
soft link



These links provide some useful explanations about ***Unix Links*** :

- <https://www.geeksforgeeks.org/soft-hard-links-unixlinux/>
- <http://linuxgazette.net/105/pitcher.html>