



**FIT9134**

# **Computer architecture and operating systems**

**Week 10**

**Operating Systems V:  
Process Management – Deadlocks and IPC**

# Deadlocks

- where processes are given exclusive access to devices, files, records etc, there is a potential for a condition known as "deadlock" to occur.
- a **Deadlock** is a situation where 2 or more processes **wait indefinitely** because the resources they need to complete are being held (& never released) by other processes.

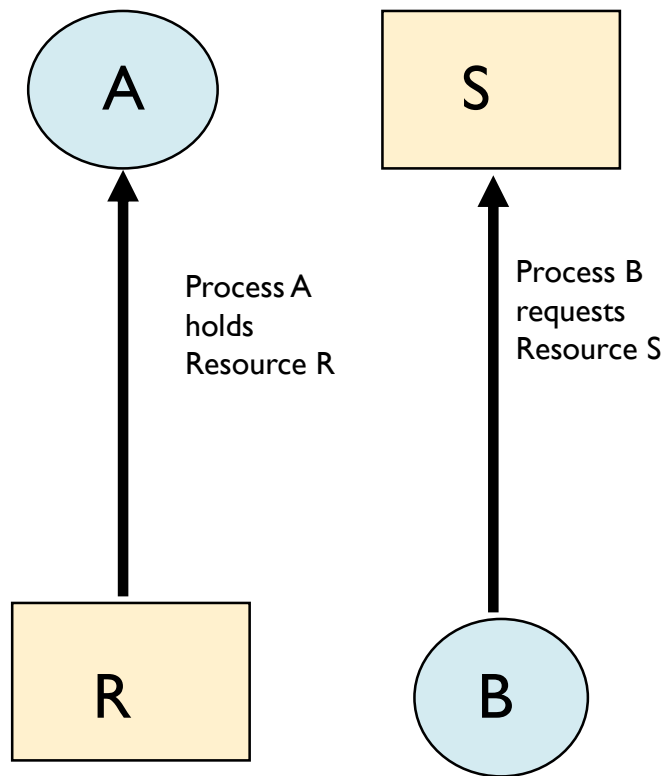
# Deadlock Conditions

4 conditions must exist simultaneously for deadlocks to occur:

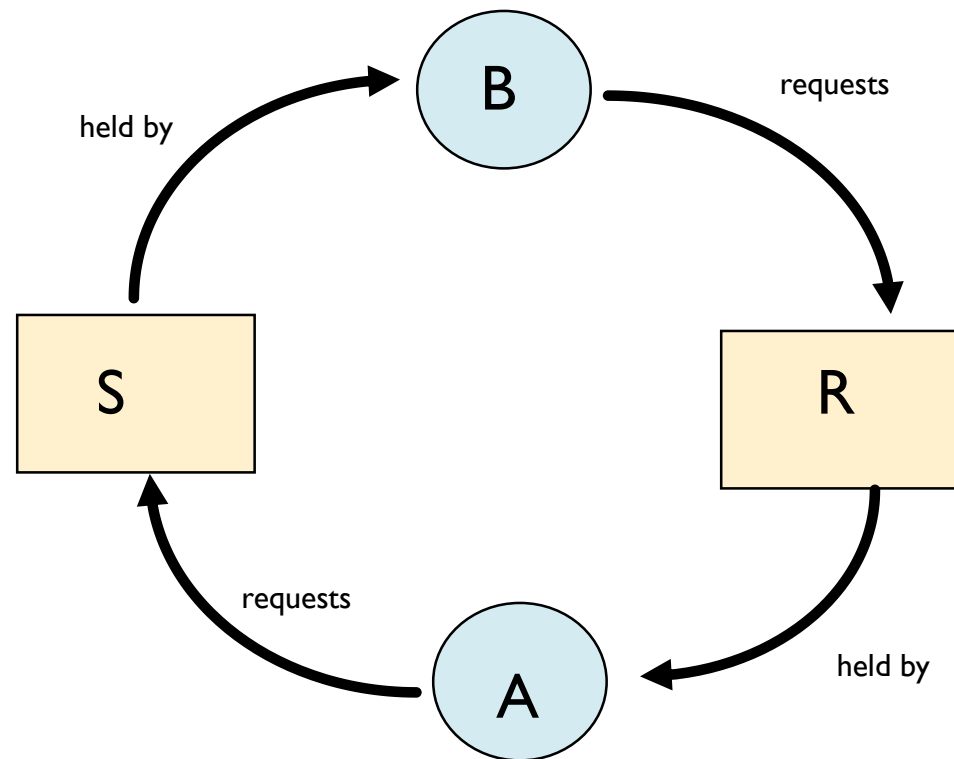
- i) **Mutual Exclusion** – a process is allowed to hold on to an un-shareable resource exclusively
- ii) **Hold & Wait** – a process is allowed to request new resources over time without releasing those it already holds
- iii) **No Pre-emption** – a resource can only be released voluntarily by the process which holds it
- iv) **Circular Wait** – 2 or more processes are waiting for resources held by each other

# Deadlock Conditions

- A **resource allocation graph** is often used to model the above conditions.



No Deadlock



Deadlock (Cycle)

# Deadlock Management

- 4 possible strategies for dealing with deadlocks:
  - i. Ostrich Algorithm
  - ii. Detection and Recovery
  - iii. Deadlock Prevention
  - iv. Deadlock Avoidance

# Deadlock Management Strategies

## i) Ostrich Algorithm

- simply ***ignore the problem***, as the overhead of dealing with all possible deadlock situations may be too high. This can sometimes be the simplest and most effective solution!

## ii) Detection and Recovery

- resource graph checked (and updated) for ***cycles*** on each request/release of resources. If cycle is about to be caused by a resource allocation, then one process is suspended, rolled back or killed.

# Deadlock Management Strategies


iii) **Deadlock Prevention** - this requires that one of four conditions as discussed before be eliminated:

- **‘mutual exclusion’** - cannot be entirely removed, since some resources are un-shareable by nature.
- **‘hold & wait’** – make all processes request resources in advance => may not be possible, and often causes non-optimal use of resources.
- **‘no pre-emption’** – forcibly removes all resources held by a process if it requests any resources which cannot be immediately allocated. Restart the process when all resource requests can be satisfied.

# Deadlock Management Strategies

- 'circular wait' – this condition is preventable if all resources are ordered/numbered, and processes are only allowed to request resources in ascending numerical order => cycles eliminated

Eg.



Deadlock will occur if A requests Q **AND** B requests P. But if P has been assigned a number higher than Q, then A will not be allowed to request Q. Similarly, if  $Q > P$ , then B is not allowed to request P. So deadlock by circular wait will not occur.



# Process Management Strategies

iv) **Deadlock Avoidance** - relies on **monitoring** the use of resources and **anticipate** requests for use.

- e.g. 'Bankers Algorithm' (refer to example in **Tanenbaum** textbook, Chapter 6), based on 'safe states'. State is safe if not deadlocked and there is a way to satisfy all future/pending requests
- Main problem with most deadlock avoidance strategies is resource requirements must be known in advance and stay static for duration of current processes. Considerable overhead involved.

# Process Coordination

- A mechanism called “**Semaphore**” may be used to provide Mutual Exclusion and Synchronization between processes
- e.g. O/S processes placing jobs in print queue and the print process removing them. The 2 processes need to be synchronized.
  - Only one process is allowed to manipulate the queue at any time; the code to do this is sometimes called a “critical code” (code which can only be executed by **one process at any one time**). We achieve the mutual exclusion and synchronization by placing the “critical code” inside a special “**Wait**” and “**Signal**” code section (using **semaphores**).

# Semaphores

- A “*semaphore*” can be thought of as a *non-negative integer* (*S*) which may only be acted upon by 2 operations:
  - *signal(S)* means :  
$$S = S + 1$$
  - *wait(S)* means :  
while **S is 0** do nothing (ie. “*waits*”)  
otherwise  $S = S - 1$
- **Signal(S)** and **Wait(S)** are **non-divisible** operations, & may only be used by one process at a time (ie. if a process wants to perform *wait(S)*, it cannot proceed until another process has performed a *signal(S)* to make the value of *S* +ve).

# Critical Section

- the “critical code to be performed” (called the ‘*critical section*’) of the process may only be executed by one process at a time.
- non-shareable resources e.g. peripherals, files, data tables can be protected from simultaneous access by processes by making those parts of code that access the resource a “**Critical Section**”. Only one process can be in a critical section at any one time.

# Critical Section example

start

createSemaphore(**add2pque**), initial value = 1

**add2pque** is a **BINARY** semaphore

....

print manager activated

verify file exists

allocate printer

add to print queue

**add2pque** becomes 0 (if no other process is in this *critical section*) and this process goes into the critical section. Otherwise this process waits.

Critical  
Section

wait(**add2pque**)

**// critical code (eg. printing) to be performed...**

signal(**add2pque**)

....

end

**add2pque** becomes 1, **signaling** that the current process is finished with the *critical section*, other processes may now enter the *critical section*.

# Semaphore Implementation

- ***wait(S)*** and ***signal(S)*** or similar inter-process communication mechanisms may be implemented as part of the instruction set of the CPU, and are used in inner kernel of the O/S.
- this allows user programs to implement critical sections in their code
- indivisibility of semaphore operation is crucial; this can be ensured by disabling interrupts on entering semaphore.

# Inter-process Communication

- As well as semaphores, Unix also supports a variety of mechanisms that processes can use to communicate with each other.
- These are referred to as **IPC (Inter-process Communication)**:
  - *Pipes*
  - *Message queues*
  - *Shared memory*

# Inter-process Communication

- These IPC (pipes, message queues & shared memory) facilities were introduced in System V Unix.
  - now used in most Unix implementations (including Linux)
- These facilities are generally accessed through system calls (calls to routines in the kernel code).



# Pipes

- A fundamental IPC facility in Unix.
- A cornerstone of the Unix philosophy of modularity.
- Pipes are used on the command line, eg:


```
$ ls -l | grep '^d' | awk '{print $4}' | sort
```

- Programs can also make use of the pipe facility using appropriate system calls.

These “Pipes” connect the **output** of one program to the **input** of another program

# Pipes

- Pipes are intended for one-way communication between processes.



- Operate on a *first-in first-out* (FIFO) basis. Data flow is one-directional.
- Flow control within the pipe is handled by the kernel on behalf of processes. For example, when a pipe is full, process A will block; when the pipe is empty, process B will block.

# Message Queues

- A message is simply a sequence of characters.
- A message queue is a linked list of messages, each of a fixed maximum size.
- New messages are always added to the end of the queue. In this sense, the order of message sending is preserved.

# Message Queues

- Once a message queue has been established, processes can place messages on the queue and remove them, through system calls.
- Message queues are in a sense similar to pipes, however message queues are more versatile:
  - messages are distinct, whereas pipes just pass unformatted streams of data.
  - messages can be assigned a *type* which can be used to allow classes messages to be processed in a particular way by a single process or distribute messages to multiple processes.
  - we do not always have to read the first message in the queue first.
  - message queues are persistent

# Shared memory

- Often considered the most efficient IPC mechanism.
- Allows multiple processes (with the right permissions) to share the same memory segment.
- The shared memory segment must be created first, then attached to a process (this “attaches” the shared memory segment to the process address space). System calls are provided for control operations (eg. reads & writes) on the shared memory segment.
- Processes can “lock” the shared memory area if they have appropriate permissions.

# Some Useful Texts/Resources

- **Lister A.M. & Eager R.D.** *Fundamentals of Operating Systems, Macmillan*
- **H.M. Deitel, P. J. Deitel, D.R. Choffnes** *Operating Systems, Prentice Hall*
- **A.S.Tanenbaum** *Modern Operating Systems, Prentice Hall*
- **A.S.Tanenbaum, A.S Woodhull** *Operating Systems Design and Implementation, Prentice Hall*